

Step-by-Step Walkthrough for CS 643 Programming Assignment 2

*Cloud Computing Energy Consumption Prediction with Apache Spark
and Docker*

Kurudunje Deekshith Shetty - KS2378

May 6, 2025

GitHub Repository: [https://github.com/KDShetty11/
Energy-Consumption-Prediction-System-in-AWS-Cloud](https://github.com/KDShetty11/Energy-Consumption-Prediction-System-in-AWS-Cloud)

Docker Hub:
<https://hub.docker.com/repository/docker/kdshetty/energypred>

Contents

1	Introduction	2
2	Setting Up the AWS Cloud Environment for Parallel Training	2
2.1	Creating an EMR Cluster	2
2.2	Uploading Files to EMR Master Node	3
2.3	Copying Files to HDFS	4
3	Parallel Model Training on EMR Cluster	4
3.1	Launching Model Training	4
3.2	Monitoring Training Job	5
3.3	Saving and Downloading the Trained Model	5
4	Prediction Application on a Single EC2 Instance	6
4.1	Launching a Single EC2 Instance	6
4.2	Pre-Configuring the EC2 Instance	7
4.3	Uploading Files to EC2 Instance	8
4.4	Extracting the Model	8
4.5	Running Prediction Without Docker	8
5	Building and Deploying the Docker Container	9
5.1	Installing Docker	9
5.2	Building the Docker Image	9
5.3	Running the Docker Container	10
5.4	Pushing to Docker Hub	10
5.5	Verifying Docker Image	11
6	Repository Links	11
7	Use of ChatGPT/AI Copilots	11
8	Notes	12

1 Introduction

This document provides a comprehensive guide for completing Programming Assignment 2 in CS 643, Cloud Computing. It details the setup of an AWS cloud environment, parallel training of an energy consumption prediction model using Apache Spark on an EMR cluster, development of a prediction application on a single EC2 instance, and deployment using Docker.

Disclaimer: Although most folders or filenames read as regress or regression, the model used for this assignment is gradient-boost and not regression, the term is just a placeholder that I overlooked.

2 Setting Up the AWS Cloud Environment for Parallel Training

2.1 Creating an EMR Cluster

Step: Launch an Amazon EMR cluster with four EC2 instances for parallel model training.

- Log in to the AWS Management Console and navigate to **EMR** > **Create Cluster**.
- Configure the cluster:
 - *Release:* emr-7.8.0 (compatible with Spark 3.5.5)
 - *Applications:* Select **Spark Initiative bundle**
 - *Cluster Name:* Energy Consumption Prediction Parallel Training
 - *Cluster Configuration:* Uniform Instance
 - *bootstrap actions:* load the bootstrap.sh code

```
1 sudo pip3 install numpy pandas
```

Explanation: Ensures NumPy and Pandas are available for data processing.




Figure 1: Python Dependency Bootstrap

- Hardware configuration:
 - *Instance Type*(both primary and core: m5.xlarge
 - *Number of Instances:* 4 (1 primary, 3 Core instances)
- Security: Select an EC2 key pair (vokey.ppk).
- IAM: Choose EMRDefault (add a rule to enable ssh).
- Click **Create Cluster**.

Explanation: This sets up a managed Spark cluster for distributed training on AWS EMR.

Instance groups (2) [Info](#) ⌂ Terminate instance Resize instance group Add task instance group

With the instance groups configuration, each node type consists of the same instance type and the same purchasing option for instances: On-Demand or Spot.

Find instances by status ▼ Q Find resources by ID or type; or search for text within loaded results < 1 > ⚙

Type and name	ID	Status Last state change reason	Instances	Purchasing option and p...	EBS size (GiB)
<input type="radio"/> Primary	ig-25CS6ZVODZ1V1	Running	1	On-Demand	-
<input type="radio"/> Core	ig-2YERBRIHFMGNZ	Running	3	On-Demand	-

Figure 2: Screenshot of EMR Cluster Running

2.2 Uploading Files to EMR Master Node

Step: Transfer datasets and training script to the EMR master node using SFTP.

- Wait for the EMR cluster to reach the **Waiting** state.
- Copy the master node's public DNS (hadoop@ec2-3-83-174-73.compute-1.amazonaws.com).
- Open a terminal and start an SFTP session:

```
1 sftp -i labuser.pem hadoop@ec2-3-83-174-73.compute-1.amazonaws.com
```

- Upload files:

```
1 put TrainingDataset.csv
2 put train_model.py
```

Explanation: Transfers necessary files to the master node for storage in HDFS. Alternatively you can use winscp as shown below in the figure




Figure 3: SFTP File Upload to EMR Master Node

2.3 Copying Files to HDFS

Step: Access the master node via SSH and move files to HDFS.

- SSH into the master node(Alternatively can use putty):

```
1 ssh -i labuser.pem hadoop@ec2-3-83-174-73.compute-1.amazonaws.com
```

- Copy files to HDFS:

```
1 hadoop fs -put TrainingDataset.csv /user/hadoop/TrainingDataset.csv
2 hadoop fs -put train_model.py /user/hadoop/train_model.py
```

- Verify files:

```
1 hdfs dfs -ls -t -R
```

Explanation: Stores datasets and script in HDFS for distributed access by Spark.

```
Using username "hadoop".
Authenticating with public key "imported-openssh-key"

#_
~\#### Amazon Linux 2023
~\_#####
~\_#####
~\_###|
~\_#/ https://aws.amazon.com/linux/amazon-linux-2023
~V~'-'>
~~
~~~
~~~
~/m/'

EEEEEEEEEEEEEEEEEEEE MMMMMMM      MMMMMMM RRRRRRRRRRRRRR
E::::::::::::::::::E M:::M          M:::M R:::::::::R
EE:::EEEEEEE:::E M:::M          M:::M R:::RRRRR:::R
 E:::E        EEEE M:::M          M:::M RR:::R    R:::R
 E:::E        M:::M M:::M M:::M M:::M R:::R    R:::R
 E:::EEEEEEEE M:::M M:::M M:::M M:::M R:::RRRRR:::R
 E:::EEEEEEEE M:::M M:::M M:::M M:::M R:::RRRRR:::R
 E:::E        M:::M M:::M M:::M M:::M R:::R    R:::R
 E:::E        EEEE M:::M   MM M:::M M:::M R:::R    R:::R
 EE:::EEEEEEE:::E M:::M          M:::M R:::R    R:::R
 E::::::::::::::::::E M:::M          M:::M RR:::R  R:::R
 EEEEEEEEEEEEEEEEE MMMMMMM      MMMMMMM RRRRRRR      RRRRRR

[hadoop@ip-172-31-81-133 ~]$ ls
TrainingDataset.csv train_model.py
[hadoop@ip-172-31-81-133 ~]$ hadoop fs -put TrainingDataset.csv /user/hadoop/Train
ingDataset.csv
[hadoop@ip-172-31-81-133 ~]$ hadoop fs -put train_model.py /user/hadoop/train_mo
del.py
[hadoop@ip-172-31-81-133 ~]$ hdfs dfs -ls -t -R
-rw-r--r--    1 hadoop hdfsadmingroup       35873 2025-05-05 23:17 TrainingDataset
.csv
-rw-r--r--    1 hadoop hdfsadmingroup       2825 2025-05-05 23:17 train_model.py
[hadoop@ip-172-31-81-133 ~]$
```

Figure 4: HDFS File Listing

3 Parallel Model Training on EMR Cluster

3.1 Launching Model Training

Step: Submit the training job to Spark.

```
1 spark-submit train model.py
```

Explanation: Executes the `train_model.py` script, training an ML model (e.g., linear regression) using MLlib across four EC2 instances. The model is saved to HDFS in a folder (e.g., regression).

```
[STATUS] Preparing features and labels...
[STATUS] Splitting data into training and test sets...
[STATUS] Training Gradient Boosted Trees regressor...
[STATUS] Generating predictions and evaluating model...

===== Model Performance =====
Root Mean Squared Error (RMSE): 73.7177
R2 (coefficient of determination): 0.9946
=====

[STATUS] Saving the trained model to 'reg' directory...
[STATUS] Model saved successfully!
```

Figure 5: Spark-Submit Training Output

3.2 Monitoring Training Job

Step: Verify job execution via the Spark Web UI.

- Access the Spark Web UI through the EMR console's **Monitor** tab or at `http://<master-node-dns>:8080`
- Confirm job completion.
- Alternatively we can use the below command to list the HDFS to verify if the trained model is saved.

```
1 hdfs dfs -ls -t -R
```

```
hadoop@ip-172-31-81-133 ~]$ hdfs dfs -ls -t -R
-rw-r--r-- 1 hadoop hdfsadmingroup 35873 2025-05-05 23:17 TrainingDataset.csv
-rwxr-xr-x - hadoop hdfsadmingroup 0 2025-05-05 23:36 regression
-rwxr-xr-x - hadoop hdfsadmingroup 0 2025-05-05 23:36 regression/data
-rw-r--r-- 1 hadoop hdfsadmingroup 0 2025-05-05 23:36 regression/data/_SUCCESS
-rw-r--r-- 1 hadoop hdfsadmingroup 23413 2025-05-05 23:36 regression/data/part-00000-b450c8e0-e490-4b22-9afe-42298952ca6b-c000.snappy.parquet
-rwxr-xr-x - hadoop hdfsadmingroup 0 2025-05-05 23:36 regression/metadata
-rw-r--r-- 1 hadoop hdfsadmingroup 0 2025-05-05 23:36 regression/metadata/_SUCCESS
-rw-r--r-- 1 hadoop hdfsadmingroup 687 2025-05-05 23:36 regression/metadata/part-00000
-rw-r--r-- 1 hadoop hdfsadmingroup 4321 2025-05-05 23:35 train_model.py
```

Figure 6: HDFS listing Trained Model

3.3 Saving and Downloading the Trained Model

Step: Copy and compress the trained model from HDFS.

```
1 hdfs dfs -copyToLocal regression /home/hadoop/regressionmod
2 tar -czf model.tar.gz regressionmod/
```

Explanation: Retrieves and compresses the model for transfer.

Step: Download the model to your local machine via SFTP or winscp.

```
1 get regressionmod/model.tar.gz
```

Explanation: Transfers the model for use in prediction.

```
[hadoop@ip-172-31-81-133 ~]$ hdfs dfs -copyToLocal regression /home/hadoop/regressionmod
[hadoop@ip-172-31-81-133 ~]$ ls
TrainingDataset.csv  regressionmod  train_model.py
[hadoop@ip-172-31-81-133 ~]$ cd regressionmod/
[hadoop@ip-172-31-81-133 regressionmod]$ ls
data  metadata
[hadoop@ip-172-31-81-133 regressionmod]$ cd ..
[hadoop@ip-172-31-81-133 ~]$ tar -czf reg.tar.gz regressionmod/
[hadoop@ip-172-31-81-133 ~]$ ls
TrainingDataset.csv  reg.tar.gz  regressionmod  train_model.py
[hadoop@ip-172-31-81-133 ~]$
```

Figure 7: Model Copy and Compression

Name	Size	Changed	Rights	Owner
..		5/5/2025 7:07:13 PM	rw-r--r--	root
regressionmod		5/5/2025 7:39:25 PM	rw-r--r--	hadoop
reg.tar.gz	20 KB	5/5/2025 7:40:52 PM	rw-r--r--	hadoop
train_model.py	5 KB	5/5/2025 7:33:15 PM	rw-rw-r--	hadoop
TrainingDataset.csv	36 KB	5/1/2025 3:01:10 PM	rw-rw-r--	hadoop

Figure 8: SFTP Model Download

4 Prediction Application on a Single EC2 Instance

4.1 Launching a Single EC2 Instance

Step: Create an EC2 instance for prediction.

- Navigate to **EC2 > Launch Instance**.
- Select AMI: Ubuntu Server 24.04 LTS.
- Instance Type: `t2.medium`.
- Select a key pair (e.g., `vockey.pem`).
- Configure security group: Allow SSH (port 22).
- Launch the instance.

Explanation: Sets up a standalone EC2 instance for prediction.

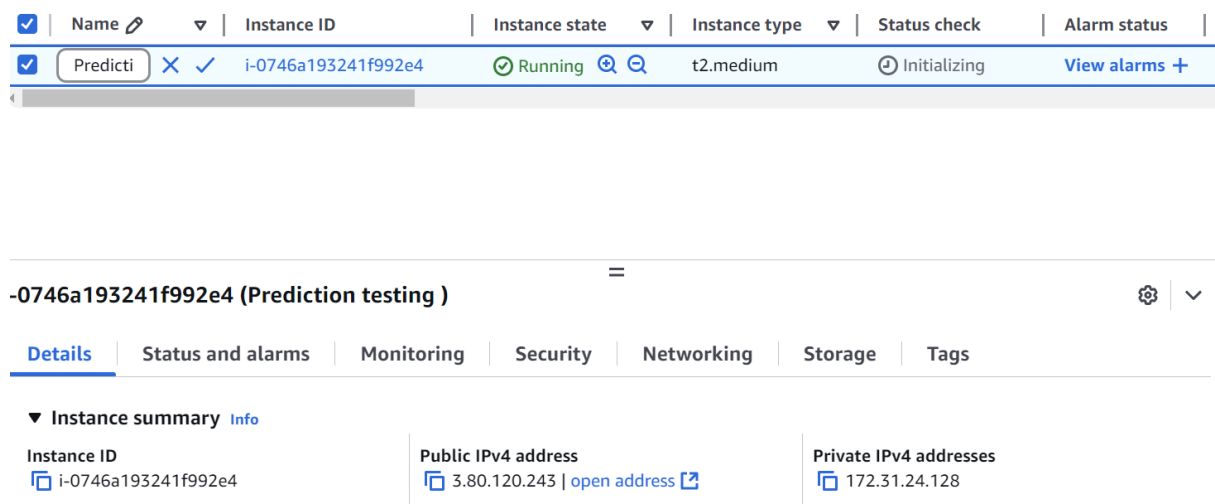


Figure 9: EC2 Instance Launch Screen

4.2 Pre-Configuring the EC2 Instance

Step: Install dependencies via SSH(can use putty).

```
1 ssh -i labuser.pem ubuntu@3.80.120.243
2 sudo apt-get update
3 sudo apt-get install -y python3-pip
4 sudo apt-get install -y python3-numpy
5 sudo apt-get install -y python3-pandas
6 sudo apt-get install -y openjdk-11-jdk
```

Explanation: Installs prerequisites for Spark.

Step: Install Apache Spark.

```
1 wget https://archive.apache.org/dist/spark/spark-3.5.5/spark-3.5.5-bin-
   hadoop3.tgz
2 sudo tar xvf spark-3.5.5-bin-hadoop3.tgz -C /opt
3 sudo chown -R ubuntu:ubuntu /opt/spark-3.5.5-bin-hadoop3
4 sudo ln -fs spark-3.5.5-bin-hadoop3 /opt/spark
```

Explanation: Configures Spark 3.5.5.

Step: Configure environment variables.

```
1 nano ~/.bash_profile
```

Add:

```
1 export SPARK_HOME=/opt/spark
2 PATH=$PATH:$SPARK_HOME/bin
3 export PATH
4 export JAVA_HOME=/usr/lib/jvm/java-1.11.0-openjdk-amd64
5 export PATH=$JAVA_HOME/bin:$PATH
```

Apply:

```
1 source ~/.bash_profile
```


Explanation: Sets up Spark and Java paths.

Step: Configure Spark logging.

```
1 cp $SPARK_HOME/conf/log4j2.properties.template $SPARK_HOME/conf/log4j2.
  properties
2 nano $SPARK_HOME/conf/log4j2.properties
```

Change `rootLogger.level = info` to `rootLogger.level = ERROR`.

Explanation: Reduces logging verbosity.

4.3 Uploading Files to EC2 Instance

Step: Upload files via SFTP.

```
1 sftp -i labuser.pem ubuntu@3.80.120.243
2 put predict.py
3 put ValidationDataset.csv
4 put model.tar.gz
```

Explanation: Transfers files for prediction. Alternatively can use winscp like the figure below.




Figure 10: SFTP Upload to EC2

4.4 Extracting the Model

Step: Extract the model archive.

```
1 tar -xzf model.tar.gz
```

Explanation: Uncompresses the model for prediction.

4.5 Running Prediction Without Docker

Step: Execute the prediction script.

```
1 spark-submit predict.py ValidationDataset.csv
```

```

ubuntu@ip-172-31-24-128:~$ ls
ValidationDataset.csv  reg.tar.gz      spark-3.5.5-bin-hadoop3.tgz
predict.py            regressionmod
ubuntu@ip-172-31-24-128:~$ spark-submit predict.py
Usage: python predict.py <test_dataset_path>
ubuntu@ip-172-31-24-128:~$

```

Figure 11: Model Extraction and Usecase

Explanation: Runs `predict.py`, which loads the model, predicts, and outputs RMSE.

Instructions for Running Without Docker:

1. Set up the EC2 instance (Sections 4.1–4.2).
2. Upload `predict.py`, `ValidationDataset.csv`, and `model.tar.gz`.
3. Extract the model: `tar -xzf model.tar.gz`.
4. Run the command, replacing `ValidationDataset.csv` with the test file path if needed.

```

ubuntu@ip-172-31-24-128:~$ ls
ValidationDataset.csv  predict.py  reg.tar.gz  regressionmod  spark-3.5.5-bin-hadoop3.tgz
ubuntu@ip-172-31-24-128:~$ spark-submit predict.py ValidationDataset.csv
[✓] Initializing SparkSession
[✓] Loading test CSV
[✓] Renaming columns
[✓] Encoding categorical columns
[✓] Dropping categorical columns
[✓] Casting columns to float
[✓] Preparing features and labels
[✓] Converting to RDD
[✓] Loading trained model
[✓] Making predictions
[✓] Pairing predictions with labels
[✓] Evaluating model

=====
          Evaluation Metrics
=====
Metric                                |   Value
-----
Root Mean Squared Error (RMSE) |  77.5037
R2 (Coefficient of Determination) |  0.9926
=====

[✓] Stopping SparkSession
ubuntu@ip-172-31-24-128:~$

```

Figure 12: Prediction Output Without Docker

5 Building and Deploying the Docker Container

5.1 Installing Docker

Step: Install Docker on the EC2 instance.

```
1 sudo apt-get install docker.io
```

Explanation: Enables Docker container operations.

5.2 Building the Docker Image

Step: Build the Docker image.

```
1 sudo docker build -t energypred .
```

Explanation: Creates an image named `energypred` using the `Dockerfile`.

5.3 Running the Docker Container

Step: Run the prediction application in a container.

```
1 sudo docker run -v /home/ubuntu/ValidationDataset.csv:/app/
  ValidationDataset.csv energypred /app/ValidationDataset.csv
```

Explanation: Maps the dataset to the container and runs the prediction.

Instructions for Running With Docker:

1. Install Docker (Section 5.1).
2. Pull the image: `docker pull kdshetty/energypred`.
3. Run the command, replacing paths as needed.

```
ubuntu@ip-172-31-24-128:~$ sudo docker run -v /home/ubuntu/ValidationDataset.csv:/app/ValidationDataset.csv energypred /app/ValidationDataset.csv
[✓] Initializing SparkSession
[✓] Loading test CSV
[✓] Renaming columns
[✓] Encoding categorical columns
[✓] Dropping categorical columns
[✓] Casting columns to float
[✓] Preparing features and labels
[✓] Converting to RDD
[✓] Loading trained model
[✓] Making predictions
[✓] Pairing predictions with labels
[✓] Evaluating model

=====
Evaluation Metrics
=====
Metric | Value
-----|-----
Root Mean Squared Error (RMSE) | 77.5037
R2 (Coefficient of Determination) | 0.9926
=====

[✓] Stopping SparkSession
ubuntu@ip-172-31-24-128:~$
```

Figure 13: Docker Container Output

```
PS D:\studies\Energy_pred\Energy-Consumption-Prediction-System-in-AWS-Cloud> docker run -v /d/studies/Energy_pred/Energy-Consumption-Prediction-System-in-AWS-Cloud/ValidationDataset.csv:/app/ValidationDataset.csv energypred /app/ValidationDataset.csv
[✓] Initializing SparkSession
[✓] Loading test CSV
[✓] Renaming columns
[✓] Encoding categorical columns
[✓] Dropping categorical columns
[✓] Casting columns to float
[✓] Preparing features and labels
[✓] Converting to RDD
[✓] Loading trained model
[✓] Making predictions
[✓] Pairing predictions with labels
[✓] Evaluating model

=====
Evaluation Metrics
=====
Metric | Value
-----|-----
Root Mean Squared Error (RMSE) | 77.5037
R2 (Coefficient of Determination) | 0.9926
=====

[✓] Stopping SparkSession
```

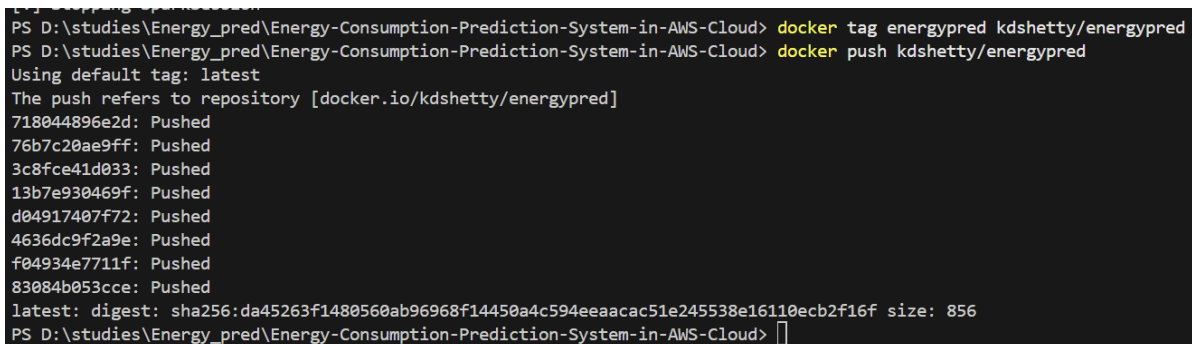
Figure 14: Docker Container Output on VSCode(Alternative)

5.4 Pushing to Docker Hub

Step: Tag and upload the image.

```
1 docker tag energypred kdshetty/energypred
2 docker push kdshetty/energypred
```

Explanation: Makes the image publicly accessible.



```
PS D:\studies\Energy_pred\Energy-Consumption-Prediction-System-in-AWS-Cloud> docker tag energypred kdshetty/energypred
PS D:\studies\Energy_pred\Energy-Consumption-Prediction-System-in-AWS-Cloud> docker push kdshetty/energypred
Using default tag: latest
The push refers to repository [docker.io/kdshetty/energypred]
718044896e2d: Pushed
76b7c20ae9ff: Pushed
3c8fce41d033: Pushed
13b7e930469f: Pushed
d04917407f72: Pushed
4636dc9f2a9e: Pushed
f04934e7711f: Pushed
83084b053cce: Pushed
latest: digest: sha256:da45263f1480560ab96968f14450a4c594eeaacac51e245538e16110ecb2f16f size: 856
PS D:\studies\Energy_pred\Energy-Consumption-Prediction-System-in-AWS-Cloud>
```

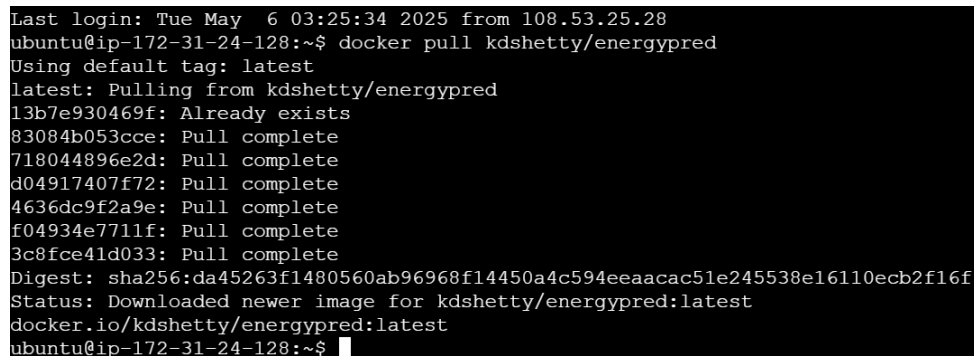
Figure 15: Docker Push to Docker Hub(Alternative on VScode)

5.5 Verifying Docker Image

Step: Pull the image to confirm availability.

```
1 docker pull kdshetty/energypred
```

Explanation: Ensures the image is deployable.



```
Last login: Tue May  6 03:25:34 2025 from 108.53.25.28
ubuntu@ip-172-31-24-128:~$ docker pull kdshetty/energypred
Using default tag: latest
latest: Pulling from kdshetty/energypred
13b7e930469f: Already exists
83084b053cce: Pull complete
718044896e2d: Pull complete
d04917407f72: Pull complete
4636dc9f2a9e: Pull complete
f04934e7711f: Pull complete
3c8fce41d033: Pull complete
Digest: sha256:da45263f1480560ab96968f14450a4c594eeaacac51e245538e16110ecb2f16f
Status: Downloaded newer image for kdshetty/energypred:latest
docker.io/kdshetty/energypred:latest
ubuntu@ip-172-31-24-128:~$
```

Figure 16: Docker Image Pull

6 Repository Links

- **GitHub Repository:** <https://github.com/KDShetty11/Energy-Consumption-Prediction-System-in-AWS-Cloud>
- **Docker Hub Repository:** <https://hub.docker.com/repository/docker/kdshetty/energypred>

7 Use of ChatGPT/AI Copilots

Code Generated by ChatGPT:

- Initial `train_model.py` structure (Spark DataFrame and MLlib Gradient Boot).
- Partial Dockerfile (base image and dependencies).

Code Written from Scratch:

- Parameter tuning in `train_model.py` for RMSE optimization.
- `predict.py` for model loading and prediction.

Code Adapted from ChatGPT:

- Modified MLlib code for dataset-specific columns and RMSE calculation.
- Adjusted Dockerfile for Spark and model inclusion.

Experience with ChatGPT:

- *Usefulness*: Accelerated setup with Spark and Docker templates. MLlib examples were mostly accurate.
- *Limitations*: Generated code used outdated APIs or incorrect paths, requiring debugging. Parameter tuning advice was generic.
- *Overall*: Effective for boilerplate but required expertise to adapt.

8 Notes

- Ensure sufficient storage and permissions for EMR and EC2. Advisable to use t2.medium or better for the EC2 Instance.
- Test prediction with `ValidationDataset.csv` to verify RMSE.
- Validate Docker container on a fresh EC2 instance.